

Introduction to the Databricks Workspace for Data Engineers

Module 01

Module Objectives

Introduction to the Databricks Workspace for Data Engineers

1. Describe the core components of the Databricks Lakehouse platform.
2. Navigate the Databricks Data Science & Engineering Workspace UI.
3. Create and manage clusters using the Databricks Clusters UI.
4. Develop and run code in multi-cell Databricks notebooks using basic operations.
5. Integrate git support using Databricks Repos.

Module Overview

Introduction to the Databricks Workspace for Data Engineers

Databricks Architecture and Services

Navigating the Databricks Workspace

Compute Resources

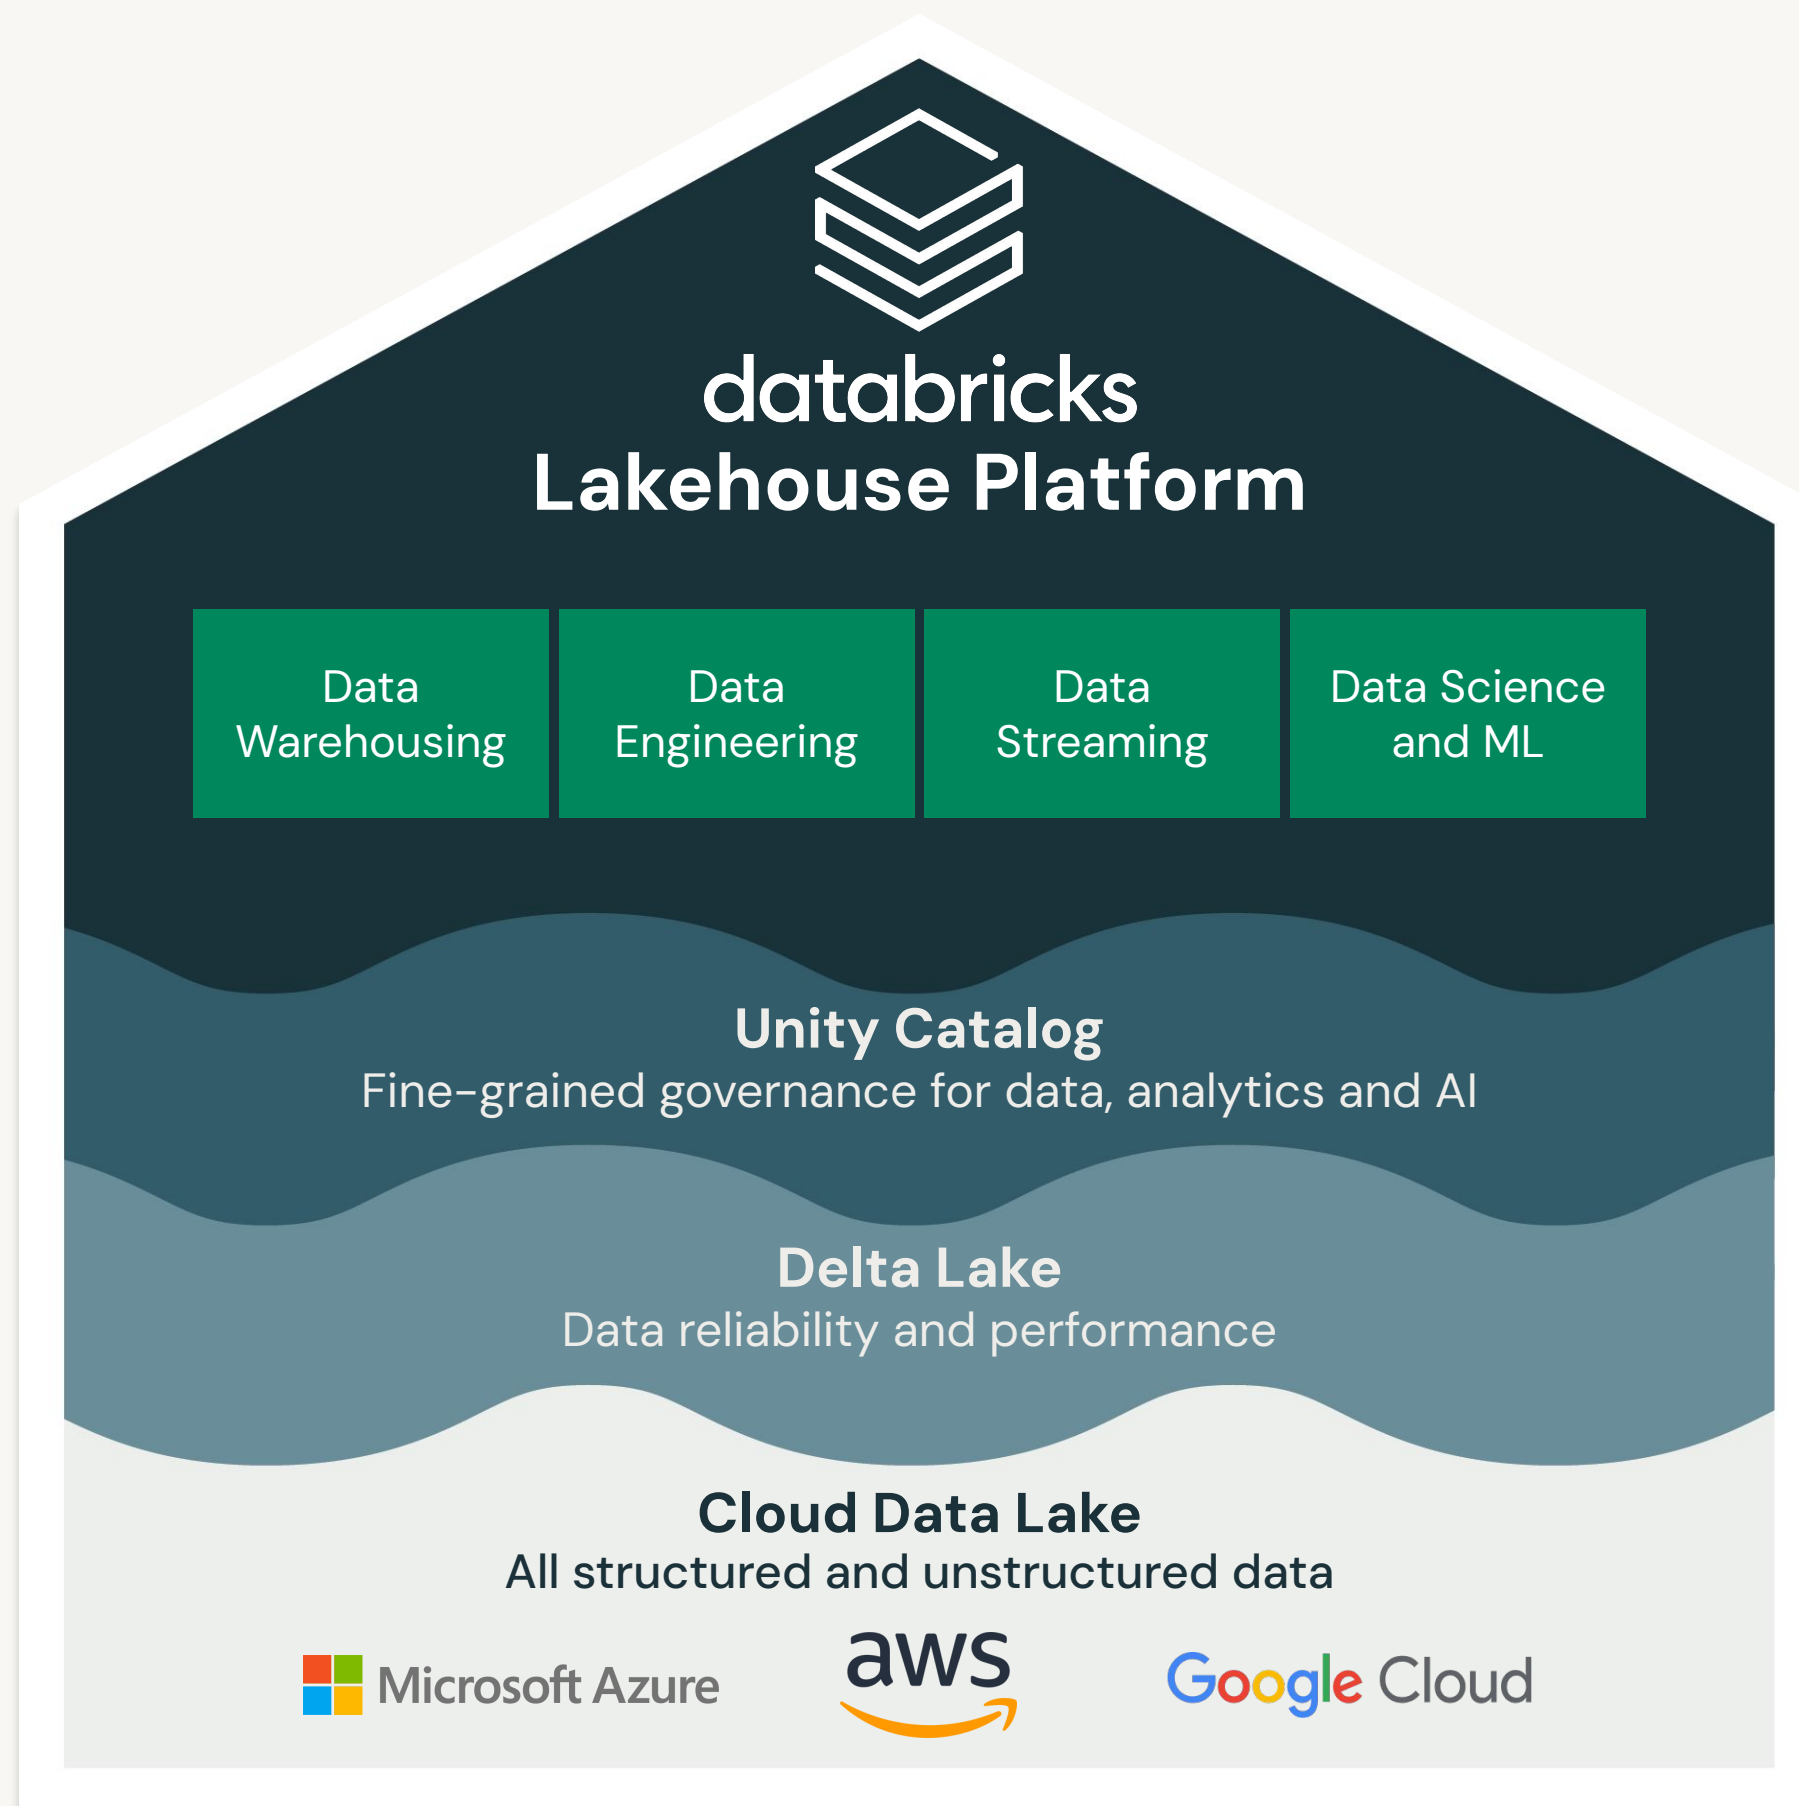
DE 1.1 – Create and Manage Interactive Clusters

Working with Databricks Repos

DE 1.2 – Databricks Notebook Operations

DE 1.3L – Get Started with the Databricks Platform Lab

Databricks Architecture and Services



Databricks Lakehouse Platform

Simple

Unify your data warehousing and AI use cases on a single platform

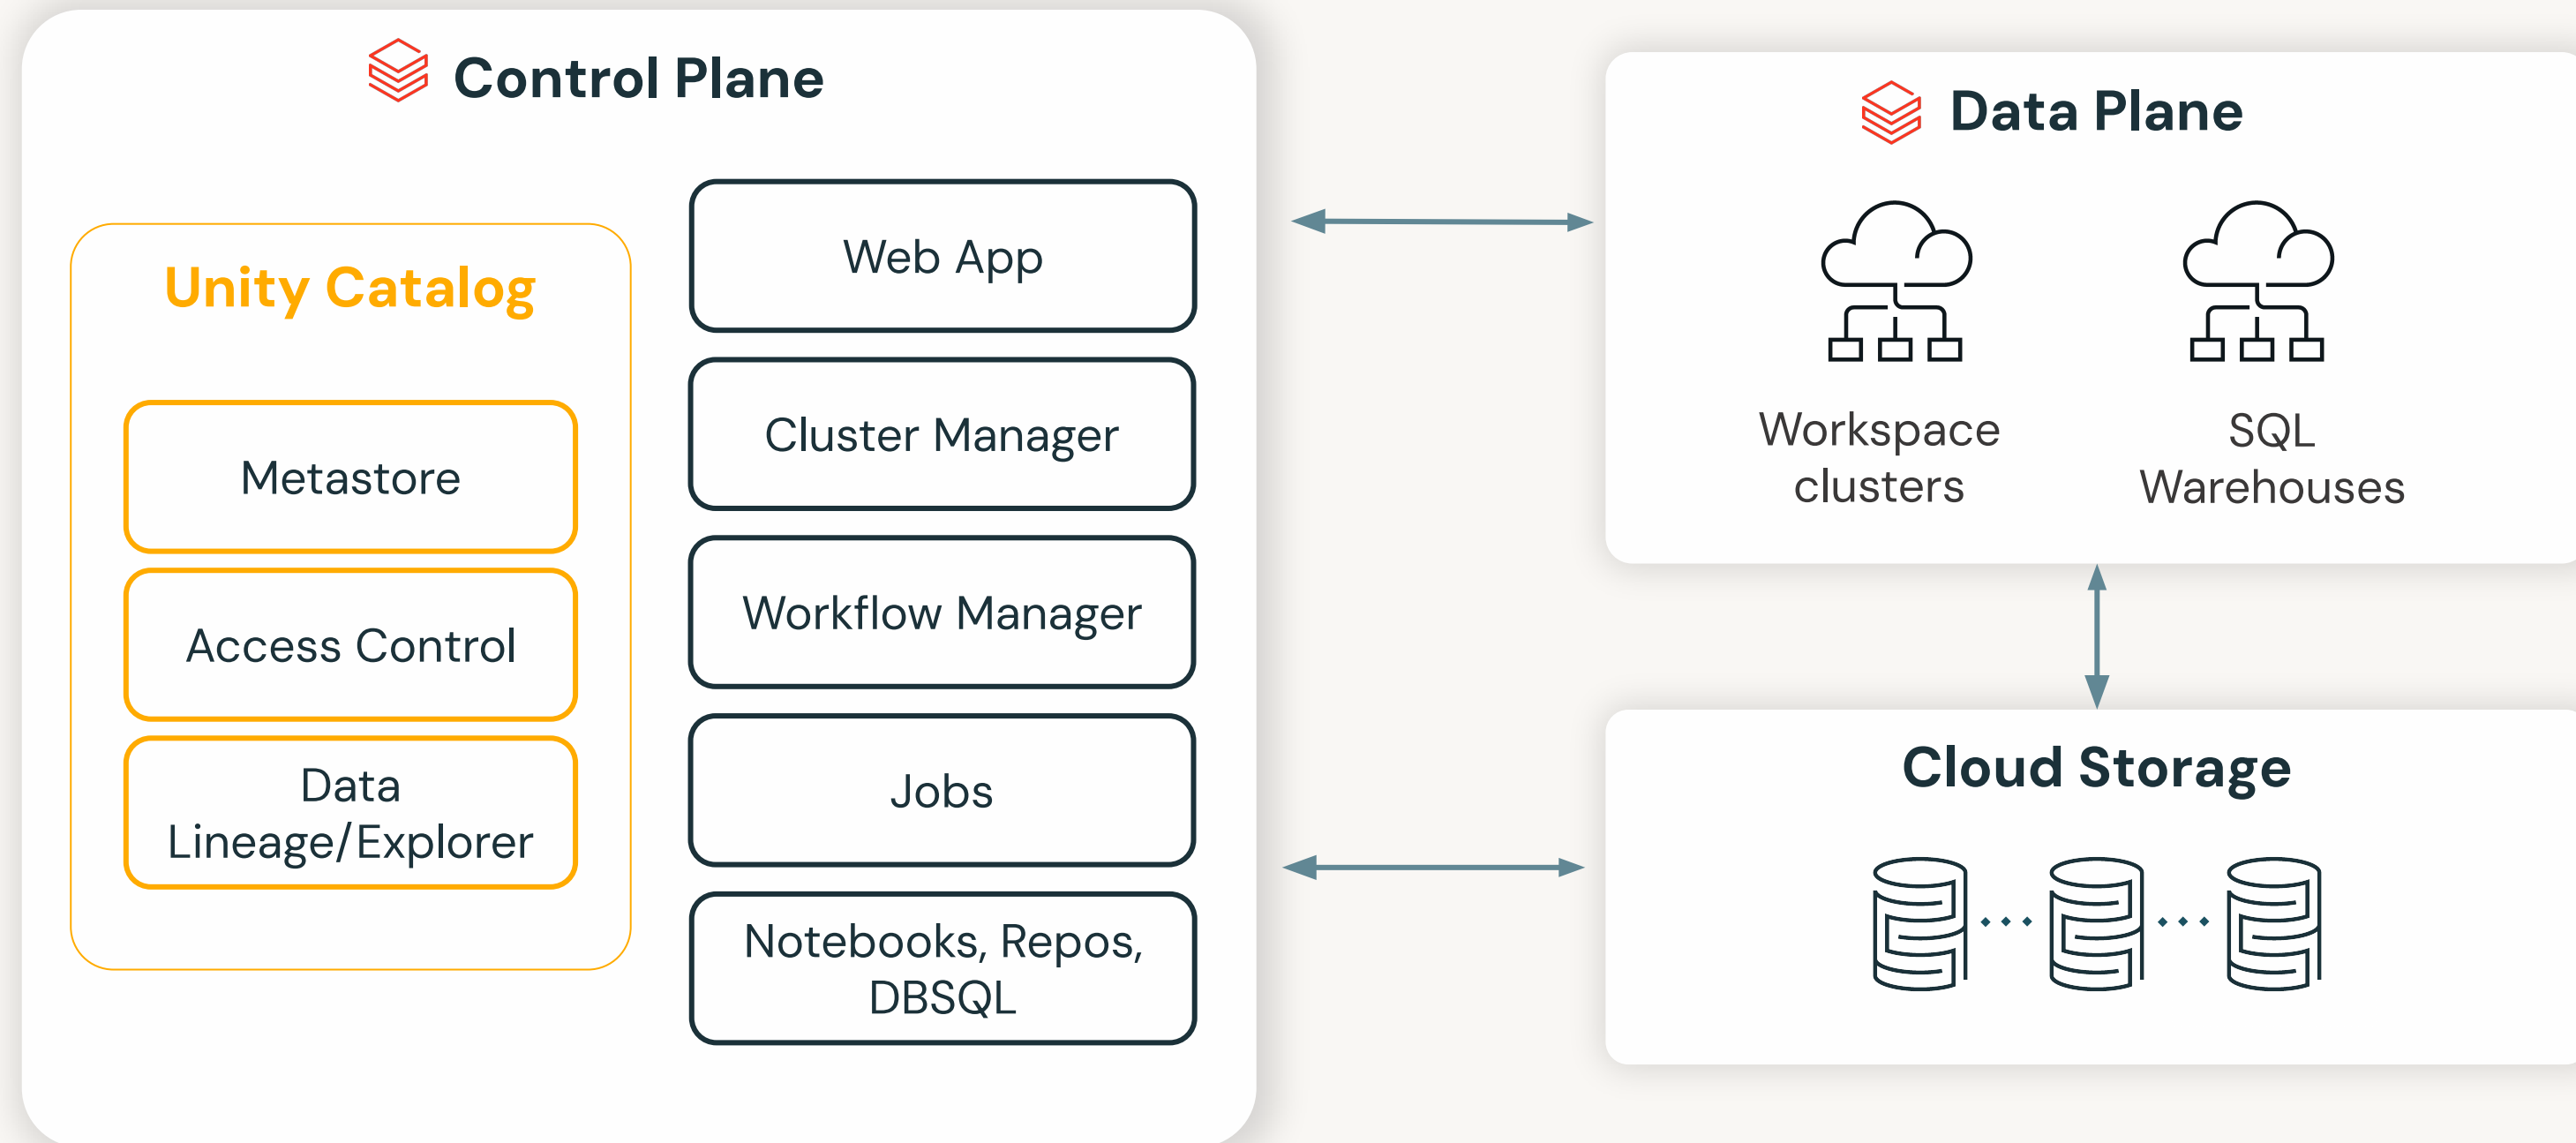
Multicloud

One consistent data platform across clouds

Open

Built on open source and open standards

Databricks Workspace and Services



Demo: Navigating the Databricks Workspace

Compute Resources

Clusters

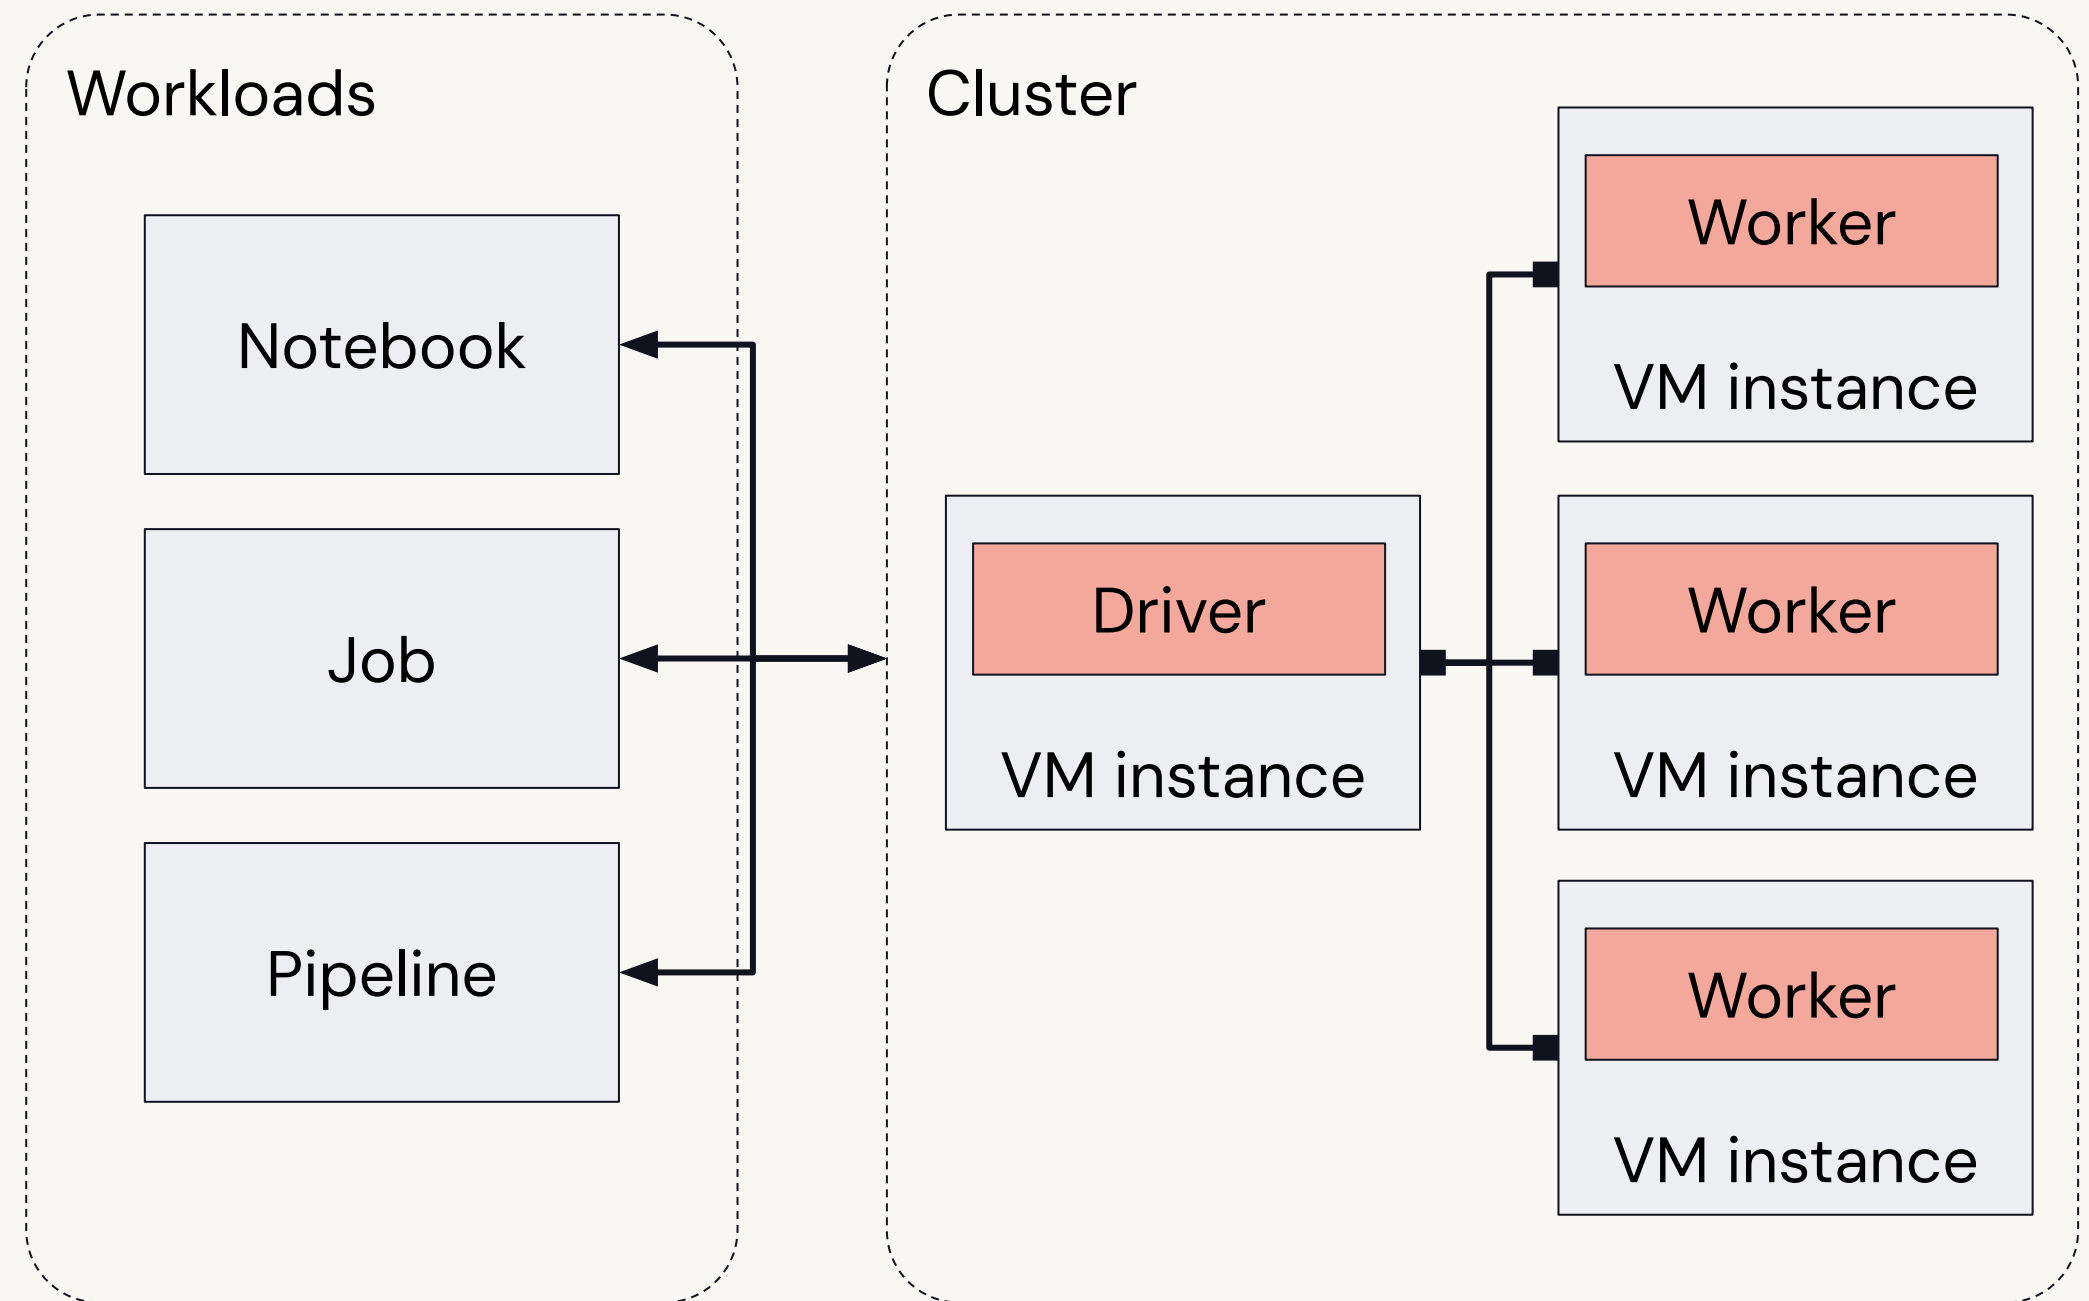
Overview

Collection of VM instances

Distributes workloads
across workers

Two main types:

1. **All-purpose** clusters for interactive development
2. **Job** clusters for automating workloads



Cluster Types

All-purpose Clusters

Analyze data collaboratively using **interactive** notebooks

Create clusters from the Workspace or API

Configuration information retained for up to 70 clusters for up to 30 days

Job Clusters

Run **automated** jobs

The Databricks job scheduler creates job clusters when running jobs

Configuration information retained for up to 30 most recently terminated clusters



Cluster Configuration

Cluster Mode

Standard (Multi Node)

Default mode for workloads developed in any supported language (requires at least two VM instances)

Single node

Low-cost single-instance cluster catering to single-node machine learning workloads and lightweight exploratory analysis

Databricks Runtime Version

Standard

Apache Spark and many other components and updates to provide an optimized big data analytics experiences

Machine learning

Adds popular machine learning libraries like TensorFlow, Keras, PyTorch, and XGBoost.

Photon

An optional add-on to optimize Spark queries (e.g. SQL, DataFrame)

Access Mode

Access mode dropdown	Visible to user	Unity Catalog support	Supported languages
Single user	Always	Yes	Python, SQL, Scala, R
Shared	Always (Premium plan required)	Yes	Python (DBR 11.1+), SQL
No isolation shared	Can be hidden by enforcing user isolation in the admin console or configuring account-level settings	No	Python, SQL, Scala, R
Custom	Only shown for existing clusters <i>without</i> access modes (i.e. legacy cluster modes, Standard or High Concurrency); not an option for creating new clusters.	No	Python, SQL, Scala, R

Cluster Policies

Cluster policies can help to achieve the following:

- Standardize cluster configurations
- Provide predefined configurations targeting specific use cases
- Simplify the user experience
- Prevent excessive use and control cost
- Enforce correct tagging

Cluster Access Control

	No Permissions	Can Attach To	Can Restart	Can Manage
Attach notebook		✓	✓	✓
View Spark UI, cluster metrics, driver logs		✓	✓	✓
Start, restart, terminate			✓	✓
Edit				✓
Attach library				✓
Resize				✓
Change permissions				✓



DE 1.1: Create and Manage Interactive Clusters

Use the Clusters UI to configure and deploy a cluster
Edit, terminate, restart, and delete clusters

Databricks Notebooks

Collaborative, reproducible, and enterprise ready

Multi-language

Use Python, SQL, Scala, and R, all in one Notebook

Collaborative

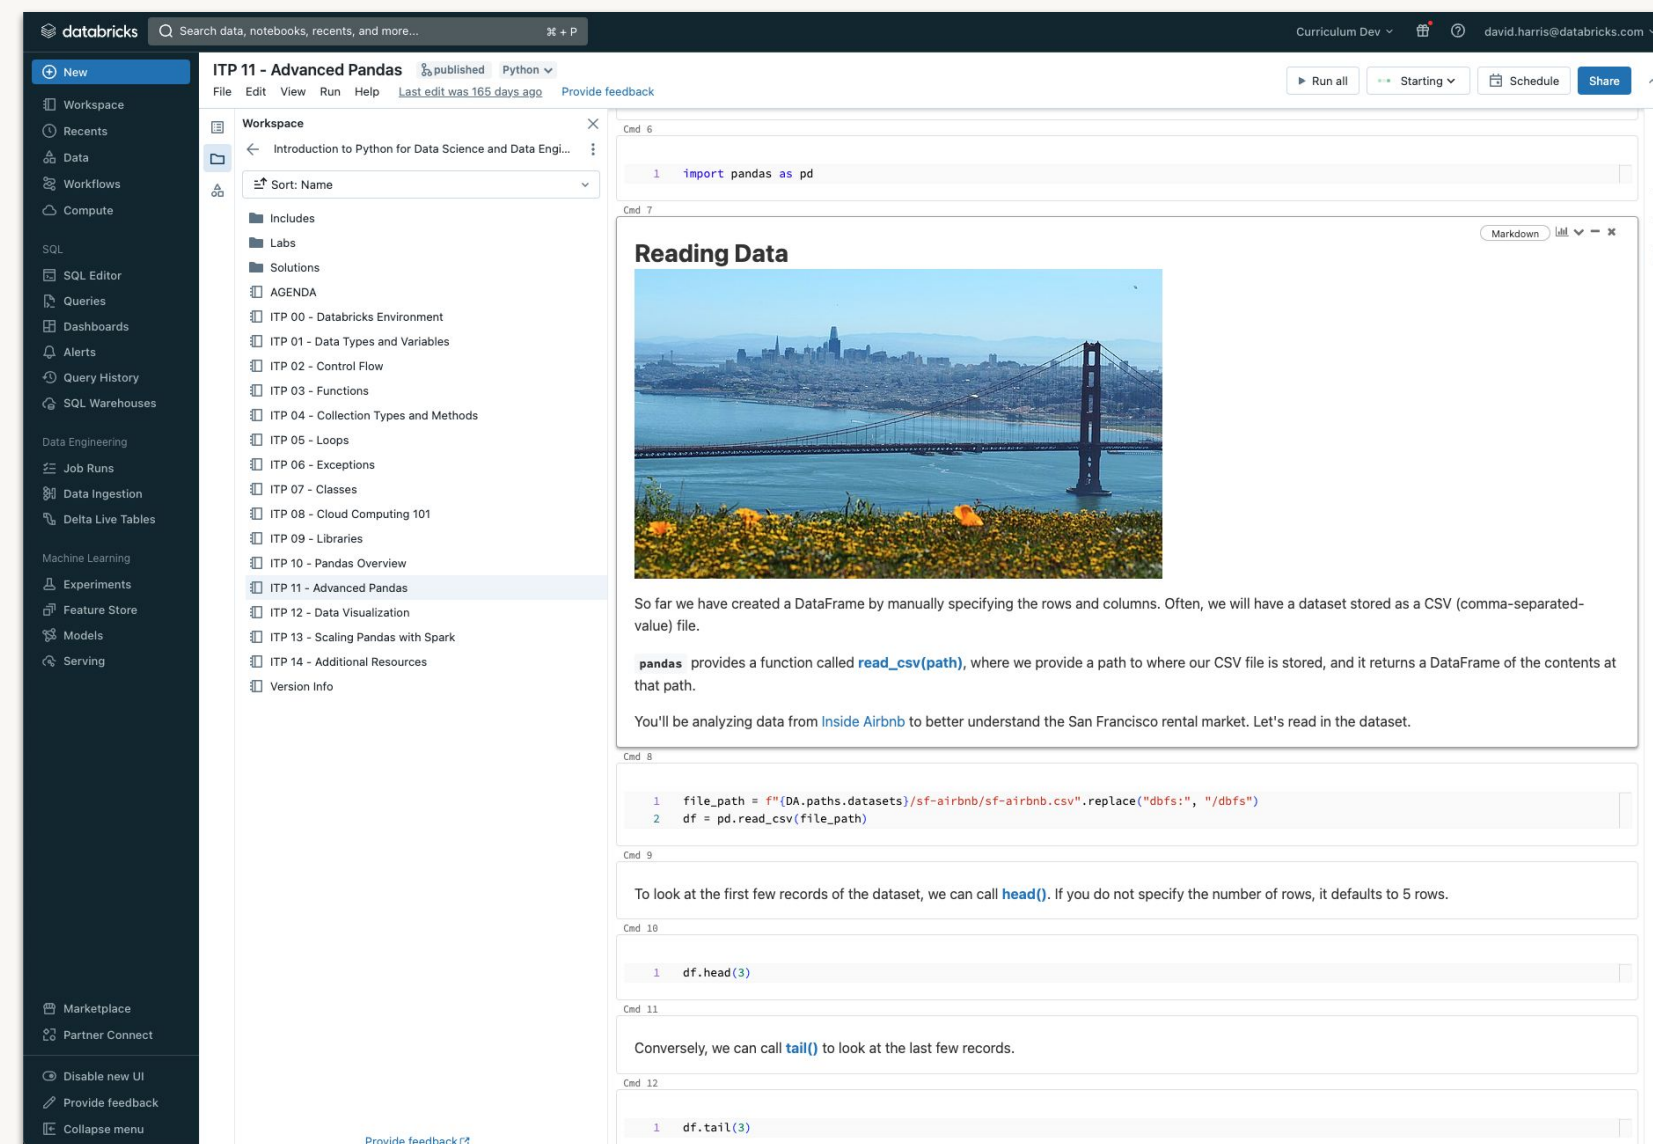
Real-time co-presence, co-editing, and commenting

Ideal for exploration

Explore, visualize, and summarize data with built-in charts and data profiles

Adaptable

Install standard libraries and use local modules



Reproducible

Automatically track version history, and use git version control with Repos

Get to production faster

Quickly schedule notebooks as jobs or create dashboards from their results, all in the Notebook

Enterprise-ready

Enterprise-grade access controls, identity management, and auditability

Notebook magic commands

Use to override default languages, run utilities/auxiliary commands, etc.

`%python, %r, %scala, %sql` Switch languages in a command cell

`%sh` Run shell code (only runs on driver node, not worker nodes)

`%fs` Shortcut for `dbutils` filesystem commands

`%md` Markdown for styling the display

`%run` Execute a remote notebook from a notebook

`%pip` Install new Python libraries

dbutils (Databricks Utilities)

Perform various tasks with Databricks using notebooks

Utility	Description	Example
<code>fs</code>	Manipulates the Databricks filesystem (DBFS) from the console	<code>dbutils.fs.ls()</code>
<code>secrets</code>	Provides utilities for leveraging secrets within notebooks	<code>dbutils.secrets.get()</code>
<code>notebook</code>	Utilities for the control flow of a notebook	<code>dbutils.notebook.run()</code>
<code>widgets</code>	Methods to create and get bound value of input widgets inside notebooks	<code>dbutils.widget.text()</code>
<code>jobs</code>	Utilities for leveraging jobs features	<code>dbutils.jobs.taskValues.set()</code>

Available within Python, R, or Scala notebooks

Working with Databricks Repos




Databricks Repos

Git Versioning

Native integration with
Github, Gitlab, Bitbucket
and Azure Devops


UI-based workflows



CI/CD Integration

API surface to integrate
with automation

Simplifies the
dev/staging/prod
multi-workspace story



Enterprise ready

Allow lists to avoid
exfiltration

Secret detection to avoid
leaking keys

Databricks Repos

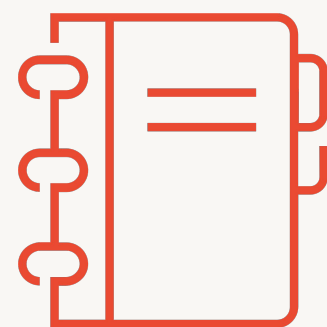
CI/CD Integration

Control Plane in Databricks

Manage customer accounts, datasets, and clusters



Databricks Web Application



Repos / Notebooks



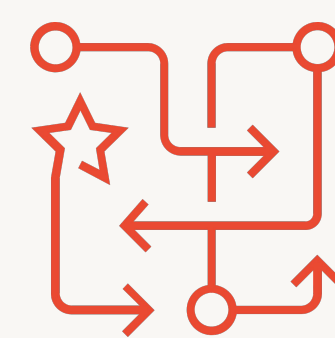
Jobs



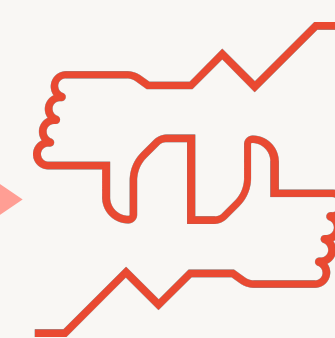
Cluster Management

Repos Service

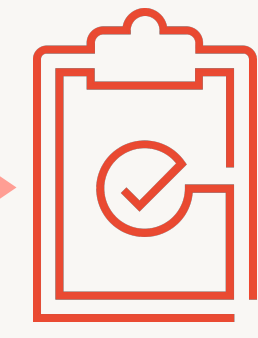
Git and CI/CD Systems




Version



Review

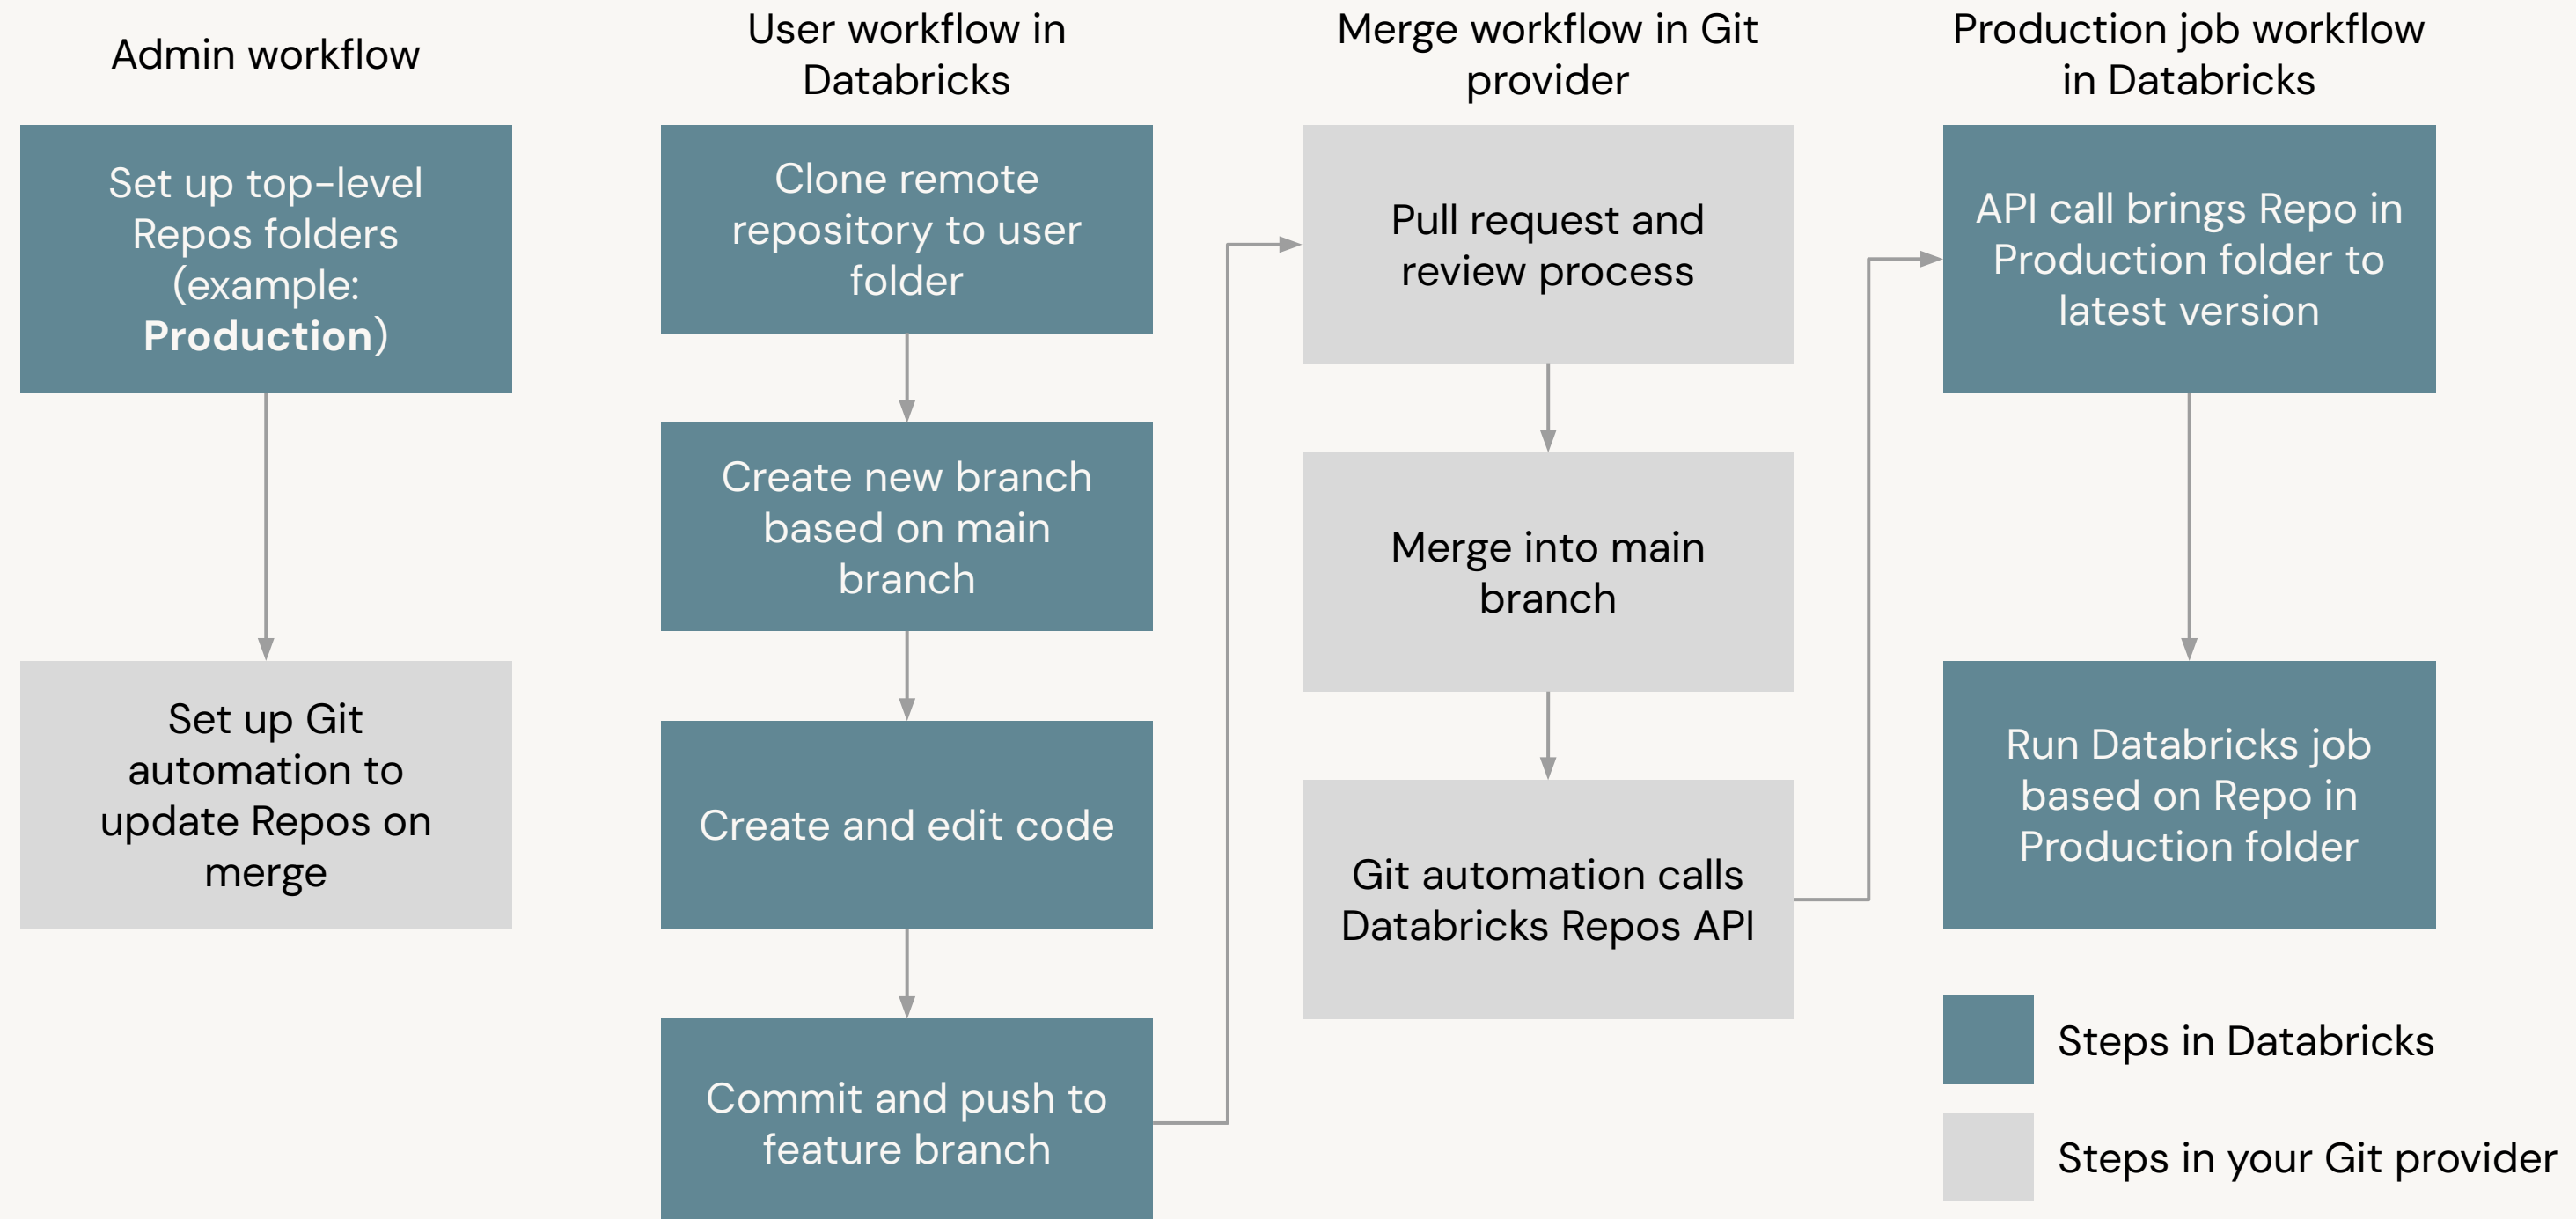


Test



CI/CD Workflow with Git and Repos

Documentation



DE 1.2: Notebook Basics

Attach a notebook to a cluster to execute a cell in a notebook

Set the default language for a notebook

Describe and use magic commands

Create and run SQL, Python, and markdown cells

Export a single or collection of notebook

DE 1.3L: Getting Started with the Databricks Platform Lab

Rename a notebook and change the default language

Attach a cluster

Use the %run magic command

Run Python and SQL cells

Create a Markdown cell



Compute Resources

Compute Resources

SQL UDFs and Control Flow

Python User-Defined Functions